

# ETC3550/ETC5550

## Applied forecasting

Ch10. Dynamic regression models

[OTexts.org/fpp3/](http://OTexts.org/fpp3/)

Rob J Hyndman  
George Athanasopoulos

# FORECASTING

## PRINCIPLES AND PRACTICE

A comprehensive introduction to the latest forecasting methods using R. Learn to improve your forecast accuracy using dozens of real data examples.



3RD EDITION

**oTexts**  
OPEN ACCESS TEXTS

# Outline

- 1 Regression with ARIMA errors
- 2 Stochastic and deterministic trends
- 3 Dynamic harmonic regression
- 4 Lagged predictors

# Outline

- 1 Regression with ARIMA errors
- 2 Stochastic and deterministic trends
- 3 Dynamic harmonic regression
- 4 Lagged predictors

# Regression with ARIMA errors

## Regression models

$$y_t = \beta_0 + \beta_1 x_{1,t} + \cdots + \beta_k x_{k,t} + \varepsilon_t,$$

- $y_t$  modeled as function of  $k$  explanatory variables  $x_{1,t}, \dots, x_{k,t}$ .
- In regression, we assume that  $\varepsilon_t$  is WN.
- Now we want to allow  $\varepsilon_t$  to be autocorrelated.

# Regression with ARIMA errors

## Regression models

$$y_t = \beta_0 + \beta_1 x_{1,t} + \cdots + \beta_k x_{k,t} + \varepsilon_t,$$

- $y_t$  modeled as function of  $k$  explanatory variables  $x_{1,t}, \dots, x_{k,t}$ .
- In regression, we assume that  $\varepsilon_t$  is WN.
- Now we want to allow  $\varepsilon_t$  to be autocorrelated.

## Example: ARIMA(1,1,1) errors

$$y_t = \beta_0 + \beta_1 x_{1,t} + \cdots + \beta_k x_{k,t} + \eta_t,$$

$$(1 - \phi_1 B)(1 - B)\eta_t = (1 + \theta_1 B)\varepsilon_t,$$

where  $\varepsilon_t$  is white noise.

# Residuals and errors

**Example:**  $\eta_t = \text{ARIMA}(1,1,1)$

$$y_t = \beta_0 + \beta_1 x_{1,t} + \cdots + \beta_k x_{k,t} + \eta_t,$$
$$(1 - \phi_1 B)(1 - B)\eta_t = (1 + \theta_1 B)\varepsilon_t,$$

# Residuals and errors

**Example:**  $\eta_t = \text{ARIMA}(1,1,1)$

$$y_t = \beta_0 + \beta_1 x_{1,t} + \cdots + \beta_k x_{k,t} + \eta_t,$$
$$(1 - \phi_1 B)(1 - B)\eta_t = (1 + \theta_1 B)\varepsilon_t,$$

- Be careful in distinguishing  $\eta_t$  from  $\varepsilon_t$ .
- Only the errors  $\varepsilon_t$  are assumed to be white noise.
- In ordinary regression,  $\eta_t$  is assumed to be white noise and so  $\eta_t = \varepsilon_t$ .

# Estimation

If we minimize  $\sum \eta_t^2$  (by using ordinary regression):

- 1 Estimated coefficients  $\hat{\beta}_0, \dots, \hat{\beta}_k$  are no longer optimal as some information ignored;
- 2 Statistical tests associated with the model (e.g., t-tests on the coefficients) are incorrect.
- 3  $p$ -values for coefficients usually too small (“spurious regression”).
- 4 AIC of fitted models misleading.



# Estimation

If we minimize  $\sum \eta_t^2$  (by using ordinary regression):

- 1 Estimated coefficients  $\hat{\beta}_0, \dots, \hat{\beta}_k$  are no longer optimal as some information ignored;
- 2 Statistical tests associated with the model (e.g., t-tests on the coefficients) are incorrect.
- 3  $p$ -values for coefficients usually too small (“spurious regression”).
- 4 AIC of fitted models misleading.
  - Minimizing  $\sum \varepsilon_t^2$  avoids these problems.
  - Maximizing likelihood similar to minimizing  $\sum \varepsilon_t^2$ .

# Regression with ARIMA errors

## Model with ARIMA(1,1,1) errors

$$y_t = \beta_0 + \beta_1 x_{1,t} + \cdots + \beta_k x_{k,t} + \eta_t,$$
$$(1 - \phi_1 B)(1 - B)\eta_t = (1 + \theta_1 B)\varepsilon_t,$$

# Regression with ARIMA errors

## Model with ARIMA(1,1,1) errors

$$y_t = \beta_0 + \beta_1 x_{1,t} + \cdots + \beta_k x_{k,t} + \eta_t,$$
$$(1 - \phi_1 B)(1 - B)\eta_t = (1 + \theta_1 B)\varepsilon_t,$$

## Equivalent to model with ARIMA(1,0,1) errors

$$y'_t = \beta_1 x'_{1,t} + \cdots + \beta_k x'_{k,t} + \eta'_t,$$
$$(1 - \phi_1 B)\eta'_t = (1 + \theta_1 B)\varepsilon_t,$$

where  $y'_t = y_t - y_{t-1}$ ,  $x'_{t,i} = x_{t,i} - x_{t-1,i}$  and  $\eta'_t = \eta_t - \eta_{t-1}$ .

## Regression with ARIMA errors

Any regression with an ARIMA error can be rewritten as a regression with an ARMA error by differencing all variables with the same differencing operator as in the ARIMA model.

# Regression with ARIMA errors

Any regression with an ARIMA error can be rewritten as a regression with an ARMA error by differencing all variables with the same differencing operator as in the ARIMA model.

## Original data

$$y_t = \beta_0 + \beta_1 x_{1,t} + \cdots + \beta_k x_{k,t} + \eta_t$$

$$\text{where } \phi(B)(1 - B)^d \eta_t = \theta(B)\varepsilon_t$$

# Regression with ARIMA errors

Any regression with an ARIMA error can be rewritten as a regression with an ARMA error by differencing all variables with the same differencing operator as in the ARIMA model.

## Original data

$$y_t = \beta_0 + \beta_1 x_{1,t} + \cdots + \beta_k x_{k,t} + \eta_t$$

$$\text{where } \phi(B)(1-B)^d \eta_t = \theta(B)\varepsilon_t$$

## After differencing all variables

$$y'_t = \beta_1 x'_{1,t} + \cdots + \beta_k x'_{k,t} + \eta'_t.$$

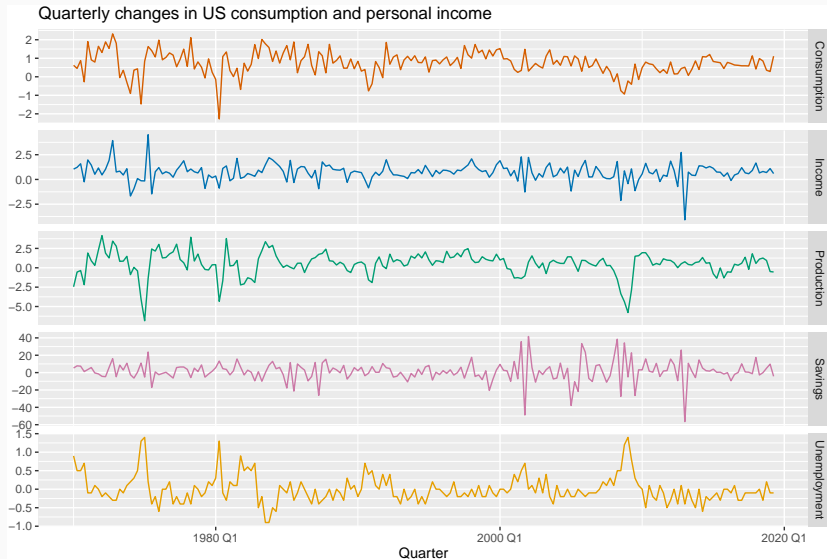
$$\text{where } \phi(B)\eta'_t = \theta(B)\varepsilon_t,$$

$$y'_t = (1-B)^d y_t, \quad x'_{i,t} = (1-B)^d x_{i,t}, \quad \text{and } \eta'_t = (1-B)^d \eta_t$$

# Regression with ARIMA errors

- In R, we can specify an ARIMA( $p, d, q$ ) for the errors, and  $d$  levels of differencing will be applied to all variables ( $y, x_{1,t}, \dots, x_{k,t}$ ) during estimation.
- Check that  $\varepsilon_t$  series looks like white noise.
- AICc can be calculated for final model.
- Repeat procedure for all subsets of predictors to be considered, and select model with lowest AICc value.

# US personal consumption and income





# US personal consumption and income

```
fit <- us_change |> model(ARIMA(Consumption ~ Income))  
report(fit)
```

```
## Series: Consumption
```

```
## Model: LM w/ ARIMA(1,0,2) errors
```

```
##
```

```
## Coefficients:
```

```
##          ar1      ma1      ma2  Income  intercept
```

```
##          0.707   -0.617   0.2066  0.1976         0.595
```

```
## s.e.    0.107    0.122   0.0741  0.0462         0.085
```

```
##
```

```
## sigma^2 estimated as 0.3113:  log likelihood=-163
```

```
## AIC=338   AICc=339   BIC=358
```

# US personal consumption and income

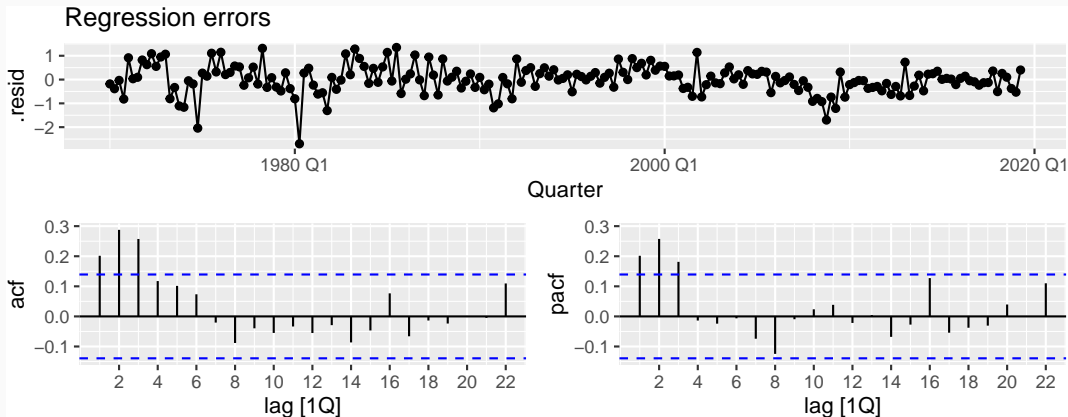
```
fit <- us_change |> model(ARIMA(Consumption ~ Income))  
report(fit)
```

```
## Series: Consumption  
## Model: LM w/ ARIMA(1,0,2) errors  
##  
## Coefficients:  
##          ar1      ma1      ma2  Income  intercept  
##          0.707  -0.617  0.2066  0.1976      0.595  
## s.e.    0.107   0.122  0.0741  0.0462      0.085  
##  
## sigma^2 estimated as 0.3113:  log likelihood=-163  
## AIC=338   AICc=339   BIC=358
```

Write down the equations for the fitted model.


# US personal consumption and income

```
residuals(fit, type = "regression") |>  
  gg_tsdisplay(.resid, plot_type = "partial") +  
  labs(title = "Regression errors")
```



# US personal consumption and income

```
residuals(fit, type = "innovation") |>  
  gg_tsdisplay(.resid, plot_type = "partial") +  
  labs(title = "ARIMA errors")
```



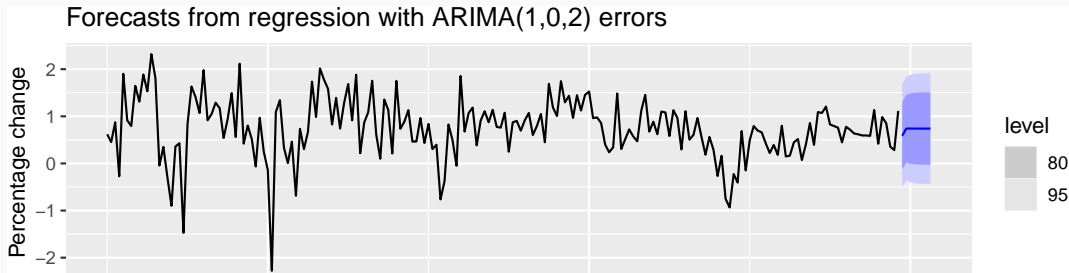
# US personal consumption and income

```
augment(fit) |>  
  features(.innov, ljung_box, dof = 5, lag = 12)
```

```
## # A tibble: 1 x 3  
##   .model                                lb_stat lb_pvalue  
##   <chr>                                <dbl>     <dbl>  
## 1 ARIMA(Consumption ~ Income)         5.54      0.595
```

# US personal consumption and income

```
us_change_future <- new_data(us_change, 8) |>
  mutate(Income = mean(us_change$Income))
forecast(fit, new_data = us_change_future) |>
  autoplot(us_change) +
  labs(
    x = "Year", y = "Percentage change",
    title = "Forecasts from regression with ARIMA(1,0,2) errors"
  )
```




# Forecasting

- To forecast a regression model with ARIMA errors, we need to forecast the regression part of the model and the ARIMA part of the model and combine the results.
- Some predictors are known into the future (e.g., time, dummies).
- Separate forecasting models may be needed for other predictors.
- Forecast intervals ignore the uncertainty in forecasting the predictors.

# Daily electricity demand

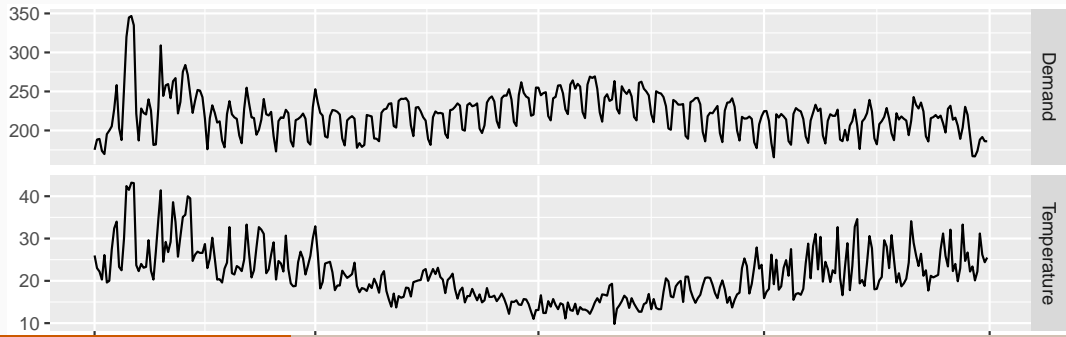
```
vic_elec_daily |>  
  ggplot(aes(x = Temperature, y = Demand, colour = Day_Type)) +  
  geom_point() +  
  labs(x = "Maximum temperature", y = "Electricity demand (GW)")
```





# Daily electricity demand

```
vic_elec_daily |>  
  pivot_longer(c(Demand, Temperature)) |>  
  ggplot(aes(x = Date, y = value)) +  
  geom_line() +  
  facet_grid(name ~ ., scales = "free_y") +  
  ylab("")
```



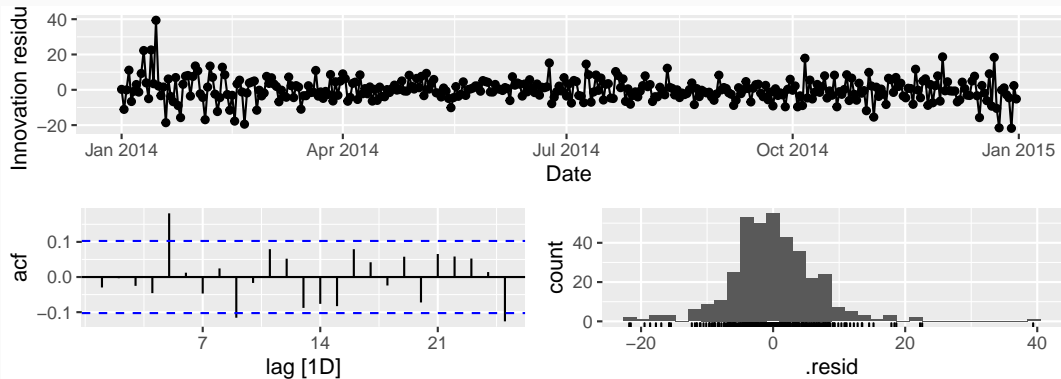
# Daily electricity demand

```
fit <- vic_elec_daily |>
  model(ARIMA(Demand ~ Temperature + I(Temperature^2) +
    (Day_Type == "Weekday")))
report(fit)
```

```
## Series: Demand
## Model: LM w/ ARIMA(2,1,2)(2,0,0)[7] errors
##
## Coefficients:
##          ar1      ar2      ma1      ma2      sar1      sar2  Temperature
##        -0.1093  0.7226 -0.0182 -0.9381  0.1958  0.417      -7.614
## s.e.      0.0779  0.0739  0.0494  0.0493  0.0525  0.057      0.448
##          I(Temperature^2)  Day_Type == "Weekday"TRUE
##                   0.1810                   30.40
## s.e.                   0.0085                   1.33
##
## sigma^2 estimated as 44.91:  log likelihood=-1206
## AIC=2432   AICc=2433   BIC=2471
```

# Daily electricity demand

```
gg_tsresiduals(fit)
```



# Daily electricity demand

```
augment(fit) |>
  features(.resid, ljung_box, dof = 9, lag = 14)

## # A tibble: 1 x 3
##   .model                                lb_stat lb_pv~1
##   <chr>                                <dbl>    <dbl>
## 1 "ARIMA(Demand ~ Temperature + I(Temperature^2) + (Day_Ty~ 28.4 3.04e-5
## # ... with abbreviated variable name 1: lb_pvalue
```

# Daily electricity demand

```
# Forecast one day ahead
```

```
vic_next_day <- new_data(vic_elec_daily, 1) |>  
  mutate(Temperature = 26, Day_Type = "Holiday")  
forecast(fit, vic_next_day)
```

```
## # A fable: 1 x 6 [1D]
```

```
## # Key:      .model [1]
```

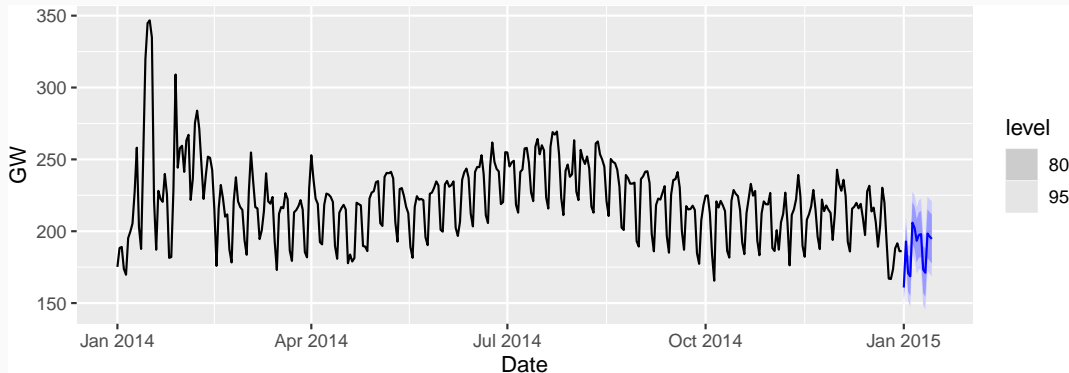
```
##   .model          Date          Demand .mean Tempe~1 Day_T~2  
##   <chr>          <date>          <dist> <dbl>   <dbl> <chr>  
## 1 "ARIMA(Demand ~ Temperature ~ 2015-01-01 N(161, 45) 161.      26 Holiday  
## # ... with abbreviated variable names 1: Temperature, 2: Day_Type
```

# Daily electricity demand

```
vic_elec_future <- new_data(vic_elec_daily, 14) |>
  mutate(
    Temperature = 26,
    Holiday = c(TRUE, rep(FALSE, 13)),
    Day_Type = case_when(
      Holiday ~ "Holiday",
      wday(Date) %in% 2:6 ~ "Weekday",
      TRUE ~ "Weekend"
    )
  )
```

# Daily electricity demand

```
forecast(fit, new_data = vic_elec_future) |>  
  autoplot(vic_elec_daily) + labs(y = "GW")
```



# Outline

- 1 Regression with ARIMA errors
- 2 Stochastic and deterministic trends
- 3 Dynamic harmonic regression
- 4 Lagged predictors



# Stochastic & deterministic trends

## Deterministic trend

$$y_t = \beta_0 + \beta_1 t + \eta_t$$

where  $\eta_t$  is ARMA process.

# Stochastic & deterministic trends

## Deterministic trend

$$y_t = \beta_0 + \beta_1 t + \eta_t$$

where  $\eta_t$  is ARMA process.

## Stochastic trend

$$y_t = \beta_0 + \beta_1 t + \eta_t$$

where  $\eta_t$  is ARIMA process with  $d \geq 1$ .

# Stochastic & deterministic trends

## Deterministic trend

$$y_t = \beta_0 + \beta_1 t + \eta_t$$

where  $\eta_t$  is ARMA process.

## Stochastic trend

$$y_t = \beta_0 + \beta_1 t + \eta_t$$

where  $\eta_t$  is ARIMA process with  $d \geq 1$ .


Difference both sides until  $\eta_t$  is stationary:

$$y'_t = \beta_1 + \eta'_t$$

where  $\eta'_t$  is ARMA process.

# Air transport passengers Australia

```
aus_airpassengers |>
  autoplot(Passengers) +
  labs(
    y = "Passengers (millions)",
    title = "Total air passengers"
  )
```



# Air transport passengers Australia

## Deterministic trend

```
fit_deterministic <- aus_airpassengers |>  
  model(ARIMA(Passengers ~ 1 + trend() + pdq(d = 0)))  
report(fit_deterministic)
```

```
## Series: Passengers  
## Model: LM w/ ARIMA(1,0,0) errors  
##  
## Coefficients:  
##          ar1  trend()  intercept  
##      0.9564    1.415    0.901  
## s.e.  0.0362    0.197    7.075  
##  
## sigma^2 estimated as 4.343:  log likelihood=-101  
## AIC=210   AICc=211   BIC=217
```

# Air transport passengers Australia

## Deterministic trend

```
fit_deterministic <- aus_airpassengers |>
  model(ARIMA(Passengers ~ 1 + trend() + pdq(d = 0)))
report(fit_deterministic)
```

```
## Series: Passengers
## Model: LM w/ ARIMA(1,0,0) errors
##
## Coefficients:
##          ar1  trend()  intercept
##      0.9564    1.415    0.901
## s.e.  0.0362    0.197    7.075
##
## sigma^2 estimated as 4.343:  log likelihood=-101
## AIC=210   AICc=211   BIC=217
```

$$y_t = 0.901 + 1.415t + \eta_t$$
$$\eta_t = 0.956\eta_{t-1} + \varepsilon_t$$
$$\varepsilon_t \sim \text{NID}(0, 4.343).$$

# Air transport passengers Australia

## Stochastic trend

```
fit_stochastic <- aus_airpassengers |>
  model(ARIMA(Passengers ~ pdq(d = 1)))
report(fit_stochastic)

## Series: Passengers
## Model: ARIMA(0,1,0) w/ drift
##
## Coefficients:
##      constant
##      1.419
## s.e.      0.301
##
## sigma^2 estimated as 4.271: log likelihood=-98.2
## AIC=200   AICc=201   BIC=204
```

# Air transport passengers Australia

## Stochastic trend

```
fit_stochastic <- aus_airpassengers |>  
  model(ARIMA(Passengers ~ pdq(d = 1)))  
report(fit_stochastic)
```

```
## Series: Passengers  
## Model: ARIMA(0,1,0) w/ drift  
##  
## Coefficients:  
##      constant  
##      1.419  
## s.e.      0.301  
##  
## sigma^2 estimated as 4.271: log likelihood=-98.2  
## AIC=200   AICc=201   BIC=204
```

$$y_t - y_{t-1} = 1.419 + \varepsilon_t,$$

$$y_t = y_0 + 1.419t + \eta_t$$

$$\eta_t = \eta_{t-1} + \varepsilon_t$$

$$\varepsilon_t \sim \text{NID}(0, 4.271).$$



# Air transport passengers Australia

```
aus_airpassengers |>
  autoplot(Passengers) +
  autolayer(fit_stochastic |> forecast(h = 20),
    colour = "#0072B2", level = 95
  ) +
  autolayer(fit_deterministic |> forecast(h = 20),
    colour = "#D55E00", alpha = 0.65, level = 95
  ) +
  labs(
    y = "Air passengers (millions)",
    title = "Forecasts from trend models"
  )
)
```



# Forecasting with trend

- Point forecasts are almost identical, but prediction intervals differ.
- Stochastic trends have much wider prediction intervals because the errors are non-stationary.
- Be careful of forecasting with deterministic trends too far ahead.

# Outline

- 1 Regression with ARIMA errors
- 2 Stochastic and deterministic trends
- 3 Dynamic harmonic regression
- 4 Lagged predictors

# Dynamic harmonic regression

## Combine Fourier terms with ARIMA errors

### Advantages


- it allows any length seasonality;
- for data with more than one seasonal period, you can include Fourier terms of different frequencies;
- the seasonal pattern is smooth for small values of  $K$  (but more wiggly seasonality can be handled by increasing  $K$ );
- the short-term dynamics are easily handled with a simple ARMA error.

### Disadvantages

- seasonality is assumed to be fixed

# Eating-out expenditure

```
aus_cafe <- aus_retail |>  
  filter(  
    Industry == "Cafes, restaurants and takeaway food services",  
    year(Month) %in% 2004:2018  
  ) |>  
  summarise(Turnover = sum(Turnover))  
aus_cafe |> autoplot(Turnover)
```




# Eating-out expenditure

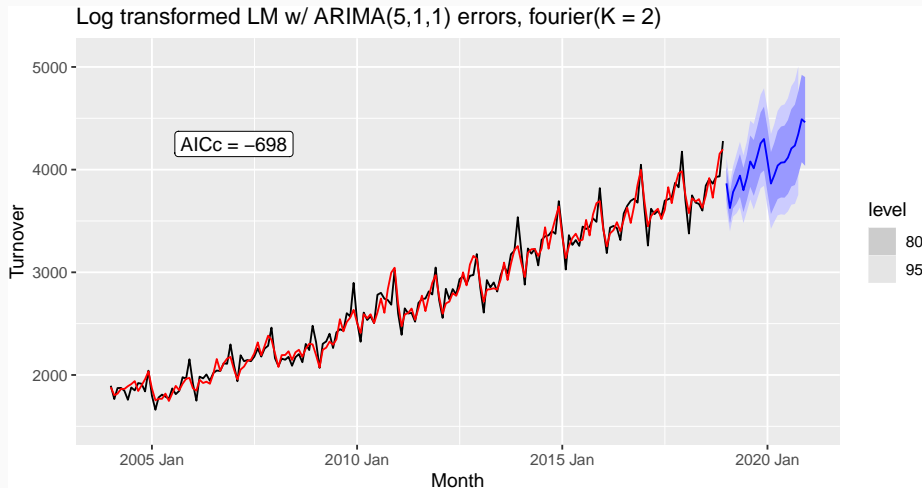
```
fit <- aus_cafe |> model(
  `K = 1` = ARIMA(log(Turnover) ~ fourier(K = 1) + PDQ(0, 0, 0)),
  `K = 2` = ARIMA(log(Turnover) ~ fourier(K = 2) + PDQ(0, 0, 0)),
  `K = 3` = ARIMA(log(Turnover) ~ fourier(K = 3) + PDQ(0, 0, 0)),
  `K = 4` = ARIMA(log(Turnover) ~ fourier(K = 4) + PDQ(0, 0, 0)),
  `K = 5` = ARIMA(log(Turnover) ~ fourier(K = 5) + PDQ(0, 0, 0)),
  `K = 6` = ARIMA(log(Turnover) ~ fourier(K = 6) + PDQ(0, 0, 0))
)
glance(fit)
```

.model	sigma2	log_lik	AIC	AICc	BIC
K = 1	0.002	317	-616	-615	-588
K = 2	0.001	362	-700	-698	-661
K = 3	0.001	394	-763	-761	-725
K = 4	0.001	427	-822	-818	-771
K = 5	0.000	474	-919	-917	-875
K = 6	0.000	474	-919	-918	-875

# Eating-out expenditure

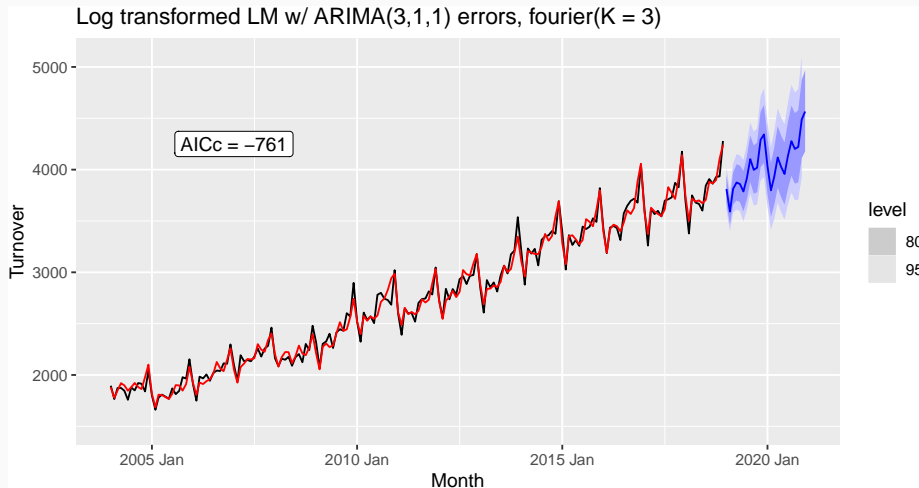


# Eating-out expenditure

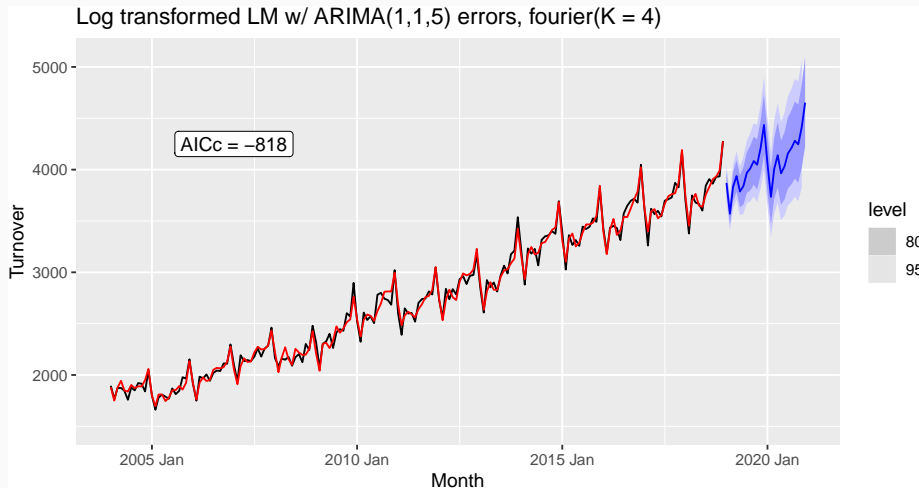





# Eating-out expenditure



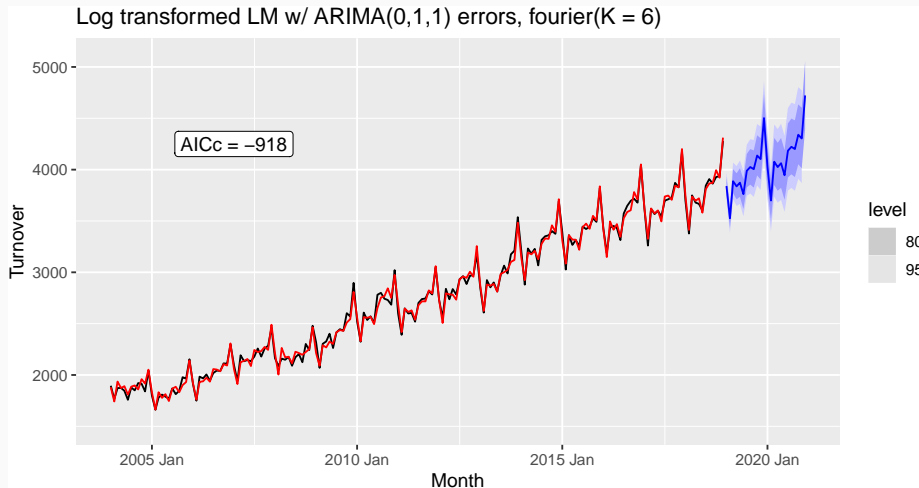
# Eating-out expenditure



# Eating-out expenditure



# Eating-out expenditure



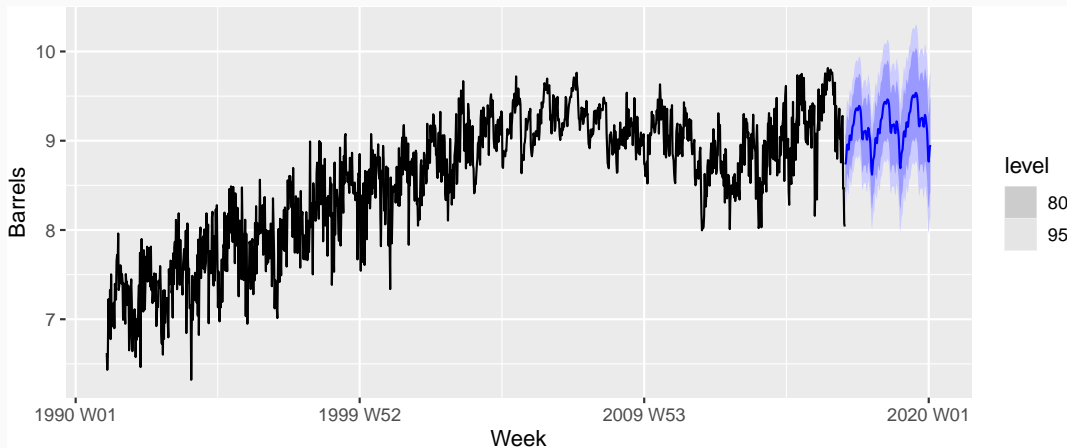
# Example: weekly gasoline products

```
fit <- us_gasoline |>
  model(ARIMA(Barrels ~ fourier(K = 13) + PDQ(0, 0, 0)))
report(fit)
```

```
## Series: Barrels
## Model: LM w/ ARIMA(0,1,1) errors
##
## Coefficients:
##          ma1  fourier(K = 13)C1_52  fourier(K = 13)S1_52
##        -0.8934             -0.1121             -0.2300
## s.e.    0.0132             0.0123             0.0122
##        fourier(K = 13)C2_52  fourier(K = 13)S2_52  fourier(K = 13)C3_52
##              0.0420             0.0317             0.0832
## s.e.         0.0099             0.0099             0.0094
##        fourier(K = 13)S3_52  fourier(K = 13)C4_52  fourier(K = 13)S4_52
##              0.0346             0.0185             0.0398
## s.e.         0.0094             0.0092             0.0092
##        fourier(K = 13)C5_52  fourier(K = 13)S5_52  fourier(K = 13)C6_52
##              -0.0315             0.0009             -0.0522
## s.e.         0.0091             0.0091             0.0090
##        fourier(K = 13)S6_52  fourier(K = 13)C7_52  fourier(K = 13)S7_52
```

# Example: weekly gasoline products

```
forecast(fit, h = "3 years") |>  
  autoplot(us_gasoline)
```



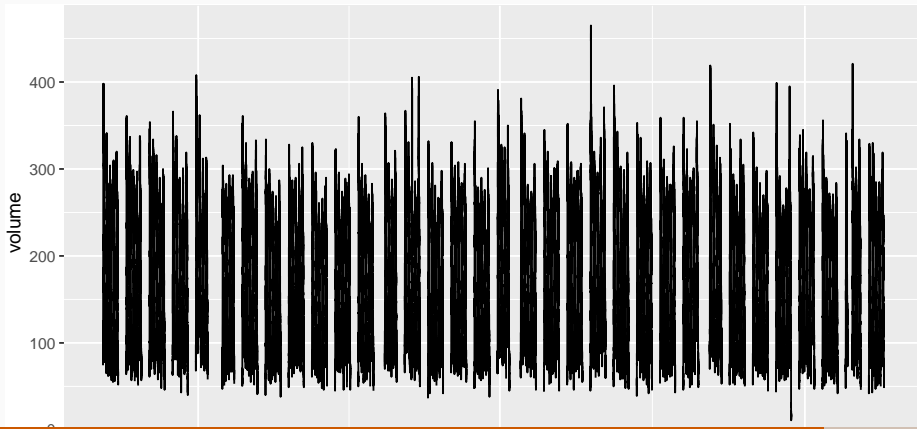
# 5-minute call centre volume

```
(calls <- readr::read_tsv("http://robjhyndman.com/data/callcenter.txt") |>
  rename(time = `...1`) |>
  pivot_longer(-time, names_to = "date", values_to = "volume") |>
  mutate(
    date = as.Date(date, format = "%d/%m/%Y"),
    datetime = as_datetime(date) + time
  ) |>
  as_tsibble(index = datetime))
```

```
## # A tsibble: 27,716 x 4 [5m] <UTC>
##   time    date      volume datetime
##   <time> <date>      <dbl> <dtm>
## 1 07:00 2003-03-03    111 2003-03-03 07:00:00
## 2 07:05 2003-03-03    113 2003-03-03 07:05:00
## 3 07:10 2003-03-03     76 2003-03-03 07:10:00
## 4 07:15 2003-03-03     82 2003-03-03 07:15:00
## 5 07:20 2003-03-03     91 2003-03-03 07:20:00
## 6 07:25 2003-03-03     87 2003-03-03 07:25:00
```

# 5-minute call centre volume

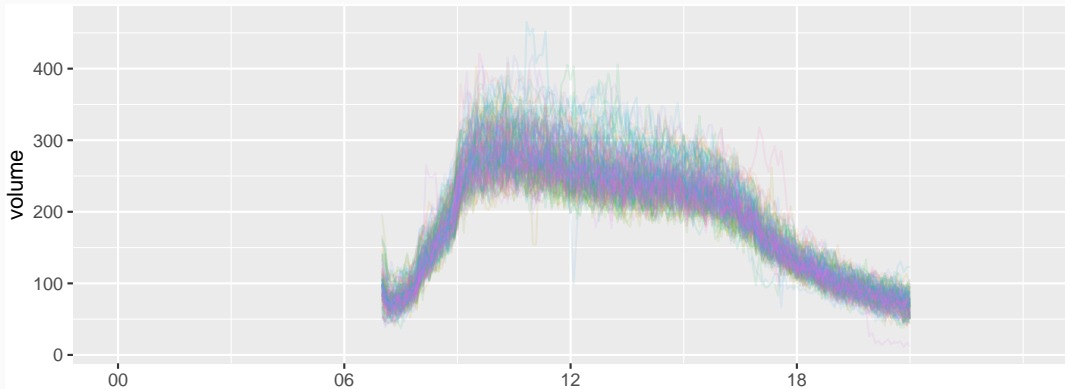
```
calls |>  
  fill_gaps() |>  
  autoplot(volume)
```





# 5-minute call centre volume

```
calls |>  
  fill_gaps() |>  
  gg_season(volume, period = "day", alpha = 0.1) +  
  guides(colour = FALSE)
```



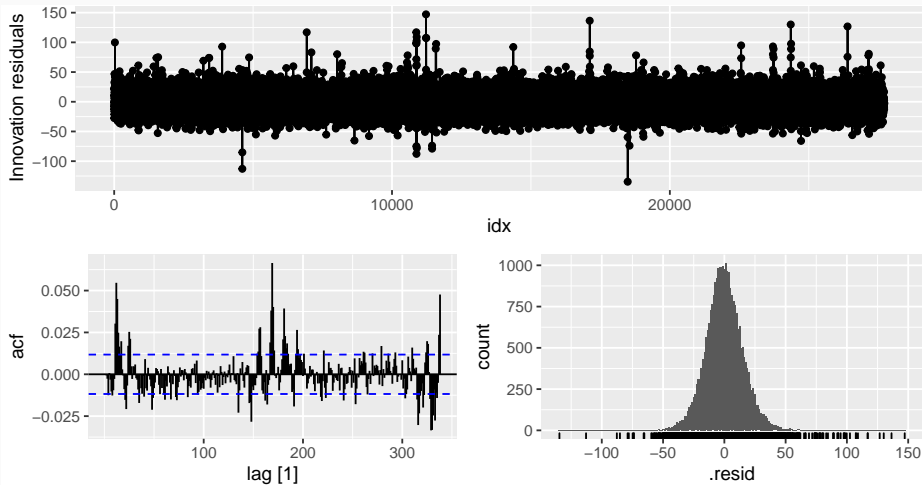
# 5-minute call centre volume

```
calls_md1 <- calls |>
  mutate(idx = row_number()) |>
  update_tsibble(index = idx)
fit <- calls_md1 |>
  model(ARIMA(volume ~ fourier(169, K = 10) + pdq(d = 0) + PDQ(0, 0, 0)))
report(fit)
```

```
## Series: volume
## Model: LM w/ ARIMA(1,0,3) errors
##
## Coefficients:
##          ar1          ma1          ma2          ma3  fourier(169, K = 10)C1_169
##          0.989    -0.7383   -0.0333   -0.0282                        -79.1
## s.e.      0.001      0.0061    0.0075    0.0060                        0.7
##          fourier(169, K = 10)S1_169  fourier(169, K = 10)C2_169
##                                55.298                        -32.361
## s.e.                                0.701                        0.378
##          fourier(169, K = 10)S2_169  fourier(169, K = 10)C3_169
```

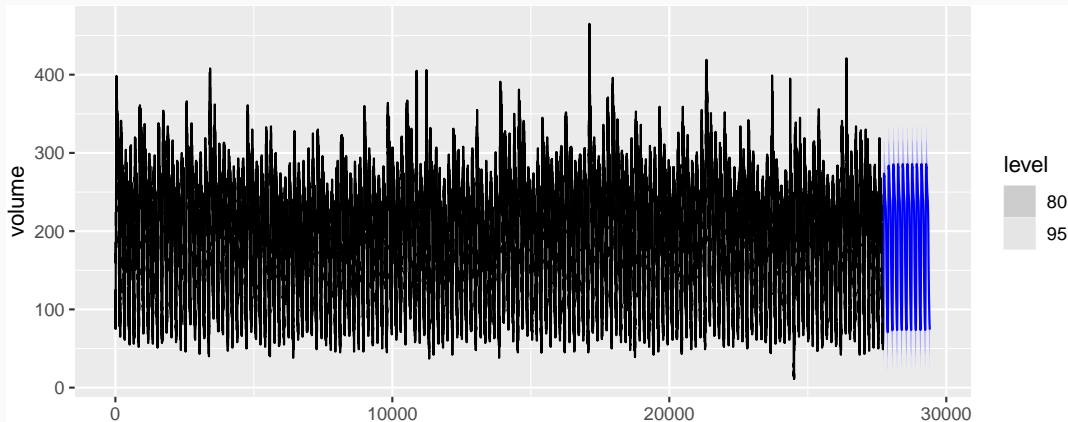
# 5-minute call centre volume

```
gg_tsresiduals(fit, lag = 338)
```



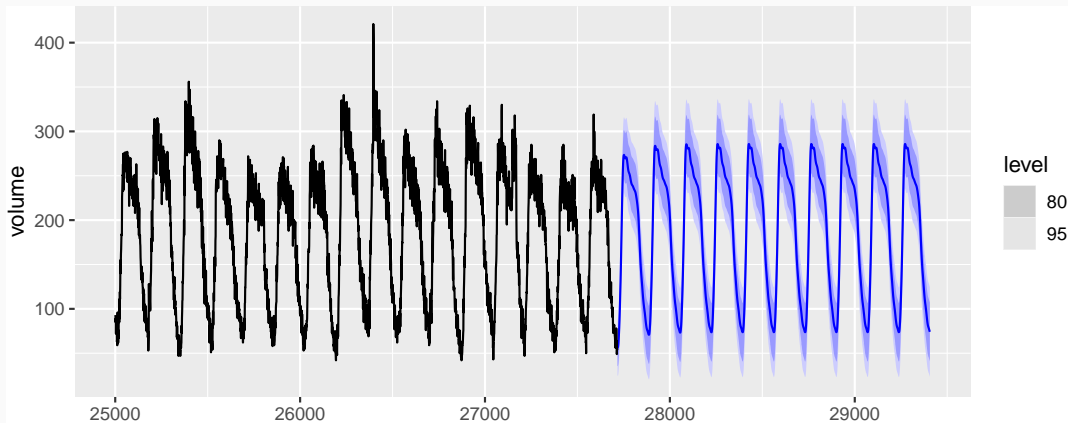
# 5-minute call centre volume

```
fit |>  
  forecast(h = 1690) |>  
  autoplot(calls_mdl)
```



# 5-minute call centre volume

```
fit |>  
  forecast(h = 1690) |>  
  autoplot(filter(calls_mdl, idx > 25000))
```



# Outline

- 1 Regression with ARIMA errors
- 2 Stochastic and deterministic trends
- 3 Dynamic harmonic regression
- 4 Lagged predictors

# Lagged predictors

Sometimes a change in  $x_t$  does not affect  $y_t$  instantaneously

# Lagged predictors

Sometimes a change in  $x_t$  does not affect  $y_t$  instantaneously

- $y_t$  = sales,  $x_t$  = advertising.
- $y_t$  = stream flow,  $x_t$  = rainfall.
- $y_t$  = size of herd,  $x_t$  = breeding stock.



# Lagged predictors

Sometimes a change in  $x_t$  does not affect  $y_t$  instantaneously

- $y_t$  = sales,  $x_t$  = advertising.
  - $y_t$  = stream flow,  $x_t$  = rainfall.
  - $y_t$  = size of herd,  $x_t$  = breeding stock.
- 
- These are dynamic systems with input ( $x_t$ ) and output ( $y_t$ ).
  - $x_t$  is often a leading indicator.
  - There can be multiple predictors.

# Lagged predictors

The model include present and past values of predictor:

$$y_t = a + \gamma_0 x_t + \gamma_1 x_{t-1} + \cdots + \gamma_k x_{t-k} + \eta_t$$

where  $\eta_t$  is an ARIMA process.

# Lagged predictors

The model include present and past values of predictor:

$$y_t = a + \gamma_0 x_t + \gamma_1 x_{t-1} + \dots + \gamma_k x_{t-k} + \eta_t$$

where  $\eta_t$  is an ARIMA process.

**Rewrite model as**

$$\begin{aligned} y_t &= a + (\gamma_0 + \gamma_1 B + \gamma_2 B^2 + \dots + \gamma_k B^k) x_t + \eta_t \\ &= a + \gamma(B) x_t + \eta_t. \end{aligned}$$

# Lagged predictors

The model include present and past values of predictor:

$$y_t = a + \gamma_0 x_t + \gamma_1 x_{t-1} + \dots + \gamma_k x_{t-k} + \eta_t$$

where  $\eta_t$  is an ARIMA process.

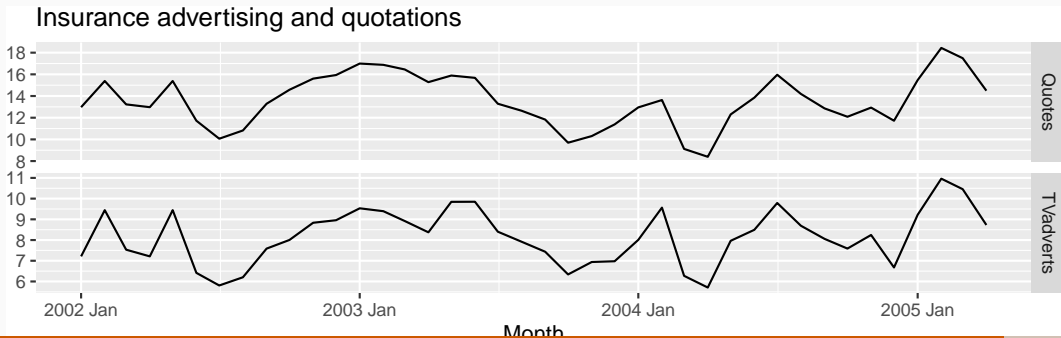
**Rewrite model as**

$$\begin{aligned} y_t &= a + (\gamma_0 + \gamma_1 B + \gamma_2 B^2 + \dots + \gamma_k B^k) x_t + \eta_t \\ &= a + \gamma(B) x_t + \eta_t. \end{aligned}$$

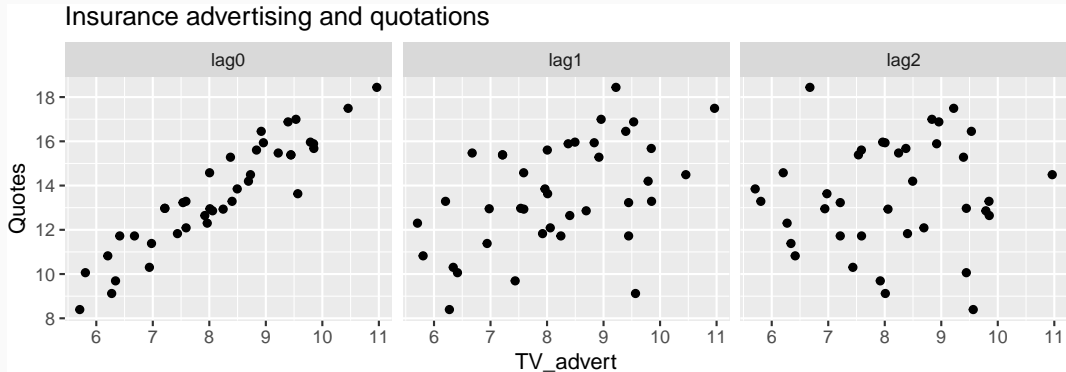
- $\gamma(B)$  is called a *transfer function* since it describes how change in  $x_t$  is transferred to  $y_t$ .
- $x$  can influence  $y$ , but  $y$  is not allowed to influence  $x$ .

# Example: Insurance quotes and TV adverts

```
insurance |>  
  pivot_longer(Quotes:TVadverts) |>  
  ggplot(aes(x = Month, y = value)) +  
  geom_line() +  
  facet_grid(vars(name), scales = "free_y") +  
  labs(y = NULL, title = "Insurance advertising and quotations")
```



# Example: Insurance quotes and TV adverts



# Example: Insurance quotes and TV adverts

```
fit <- insurance |>
  # Restrict data so models use same fitting period
  mutate(Quotes = c(NA, NA, NA, Quotes[4:40])) |>
  # Estimate models
  model(
    ARIMA(Quotes ~ pdq(d = 0) + TVadverts),
    ARIMA(Quotes ~ pdq(d = 0) + TVadverts + lag(TVadverts)),
    ARIMA(Quotes ~ pdq(d = 0) + TVadverts + lag(TVadverts) +
      lag(TVadverts, 2)),
    ARIMA(Quotes ~ pdq(d = 0) + TVadverts + lag(TVadverts) +
      lag(TVadverts, 2) + lag(TVadverts, 3))
  )
```

## Example: Insurance quotes and TV adverts

```
glance(fit)
```

Lag order	sigma2	log_lik	AIC	AICc	BIC
0	0.265	-28.3	66.6	68.3	75.0
1	0.209	-24.0	58.1	59.9	66.5
2	0.215	-24.0	60.0	62.6	70.2
3	0.206	-22.2	60.3	65.0	73.8



# Example: Insurance quotes and TV adverts

```
fit_best <- insurance |>  
  model(ARIMA(Quotes ~ pdq(d = 0) + TVadverts + lag(TVadverts)))  
report(fit_best)
```

```
## Series: Quotes
```

```
## Model: LM w/ ARIMA(1,0,2) errors
```

```
##
```

```
## Coefficients:
```

##	ar1	ma1	ma2	TVadverts	lag(TVadverts)	intercept
##	0.512	0.917	0.459	1.2527	0.1464	2.16
## s.e.	0.185	0.205	0.190	0.0588	0.0531	0.86

```
##
```

```
## sigma^2 estimated as 0.2166: log likelihood=-23.9
```

```
## AIC=61.9 AICc=65.4 BIC=73.7
```

# Example: Insurance quotes and TV adverts

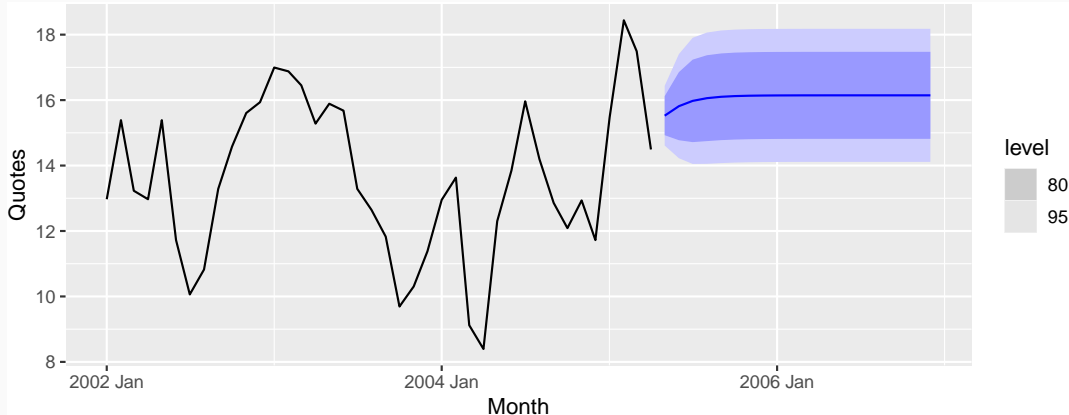
```
fit_best <- insurance |>
  model(ARIMA(Quotes ~ pdq(d = 0) + TVadverts + lag(TVadverts)))
report(fit_best)
```

```
## Series: Quotes
## Model: LM w/ ARIMA(1,0,2) errors
##
## Coefficients:
##          ar1      ma1      ma2 TVadverts lag(TVadverts) intercept
##          0.512  0.917  0.459      1.2527          0.1464         2.16
## s.e.    0.185  0.205  0.190      0.0588          0.0531         0.86
##
## sigma^2 estimated as 0.2166: log likelihood=-23.9
## AIC=61.9   AICc=65.4   BIC=73.7
```

$$y_t = 2.155 + 1.253x_t + 0.146x_{t-1} + \eta_t,$$
$$\eta_t = 0.512\eta_{t-1} + \varepsilon_t + 0.917\varepsilon_{t-1} + 0.459\varepsilon_{t-2},$$


# Example: Insurance quotes and TV adverts

```
advert_a <- new_data(insurance, 20) |>  
  mutate(TVadverts = 10)  
forecast(fit_best, advert_a) |> autoplot(insurance)
```



## Example: Insurance quotes and TV adverts

```
advert_b <- new_data(insurance, 20) |>
  mutate(TVadverts = 8)
forecast(fit_best, advert_b) |> autoplot(insurance)
```



# Example: Insurance quotes and TV adverts

```
advert_c <- new_data(insurance, 20) |>  
  mutate(TVadverts = 6)  
forecast(fit_best, advert_c) |> autoplot(insurance)
```

